CARES Reduced-Scope Plan

Plan for the ABQ demo when realistic CARES data is "a few keepers, mostly cichlids, almost nothing else." Replaces the ambitious framing in Alex-ToDo §2.1 ("HIGHEST — bulk of the work") with what's actually achievable + how to make the demo read coherently with the data we'll really have.

When: decide before §4.1 dry-run (week of 2026-05-25). Entry work itself is ~1-2 hours total.

Companion docs:

- `docs/EAZA_EEP_ENTRY_GUIDE.md` — the §2.7 walkthrough this composes with
- `docs/POPULATION_DATA_GUIDE.md` — full field-by-field walkthrough for `Institution` / `ExSituPopulation` / `HoldingRecord`
- `docs/WORKSHOP_DEMO_RUNBOOK.md` — the actual demo sequence + recovery paths

---

0. Reality check

What was assumed in the original §2.1 framing:

- Tens of CARES keepers across multiple Madagascar species families
- Bulk-load via `seed_populations` from a CSV
- Panel 2 (Studbook status) populated across all four buckets
- "Comprehensive coordinator triage view" demo narrative

What is actually true:

- A handful of CARES keepers (single digits)
- A small set of cichlid species — `Paretroplus`, possibly some `Ptychochromis`
- Almost no captive populations for other endemic families (Bedotiidae, Aplocheilidae, Anchariidae, gobies, cave fish — all effectively zero on the CARES side)
- Manual admin entry is more efficient than the CSV path for this volume; `seed_populations` is overkill

This is not a failure mode. It's the actual conservation gap the platform exists to surface. **Lean into it.**

---

1. The narrative pivot

Two demo framings are available given the data:

Framing	What it says	When it works
"Coordinator triage view" *(original §2.1)*	"Here is the state of ex-situ coordination across Madagascar endemics."	Dense data across all four Panel-2 buckets. **Not us in 2026-06.**
"The gap is the story"	"Here is what we can see, what we can't, and where the conservation gap really sits. The empty buckets are the platform's most honest output."	Sparse data, broad coverage gap. **Us.**

The "gap is the story" framing is **stronger** for an ABQ audience — SHOAL, AZA / EAZA partners, CARES leads. They already know the gap exists; what's new is having a public registry that quantifies it.

The Panel 1 hero stat (*"N of M threatened species have no known captive population"*) does most of the work. Panel 2 supports it. The other panels round out the platform-shape demonstration.

---

2. Minimum viable real-data entries

Three tiers. Do tier 1 always; tier 2 if you have time; tier 3 only if you genuinely have the data.

Tier 1 — anchor the EEPs (15 minutes)

Required for the EAZA EEP entries from §2.7 to register as actual captive populations on the dashboard. Without this, §2.7 alone leaves Panel 2 still empty.

Species	Institution	`studbook_managed`
`Bedotia madagascariensis`	Bristol Zoo Project	✓
`Paretroplus menarambo`	National Aquarium Denmark	✓

Enter these via §5 of `EAZA_EEP_ENTRY_GUIDE.md`. Use `count_total = 0` (or your best-known number) — the placeholder is fine. The entry exists primarily so Panel 2 reads "Studbook-managed: 2" rather than zero.

Tier 2 — your real CARES keepers

For each CARES keeper you actually have data for, follow the manual admin path in `POPULATION_DATA_GUIDE.md` §3 (Institution) + §4 (ExSituPopulation). Realistic scope:

- Maybe 3-5 `Institution` rows of `institution_type=hobbyist_keeper`
- Maybe 5-10 `ExSituPopulation` rows, mostly cichlids
- `studbook_managed = unchecked` (CARES is hobbyist-coordinated, not studbook-coordinated by default)
- `breeding_status = breeding` if the keeper confirms active spawning, else `unknown`

This is the data you actually have. Entering it puts cichlid species into Panel 2's "Breeding, not studbook" or "Holdings only" buckets, which is correct + tells a real story ("hobbyist keepers are doing the work the zoos aren't").

Tier 3 — known-public zoo holdings (optional)

If you want to round out the demo with a few more institutions without inventing data, the following are **public knowledge** as of 2026 (cite `April_2026_8e69dc12b4.pdf` or institutional websites):

- Toyota Aquarium (Japan) — Madagascar cichlid holdings, public
- Vienna Zoo (Tiergarten Schönbrunn) — Madagascar fish holdings, public
- Berlin Aquarium (Zoo Berlin) — Madagascar cichlid holdings, public

Only add these if you can find a source you'd cite at ABQ. Don't guess. The "gap is the story" narrative is undermined by inflated data more than by sparse data.

---

3. Interaction with the synthetic demo layer

`backend/populations/management/commands/seed_demo_coordination.py` generates synthetic `CoordinatedProgram` / `Transfer` / `BreedingRecommendation` / `BreedingEvent` rows tagged with

the `[DEMO_SEED]` marker. The `WORKSHOP_DEMO_RUNBOOK` runs it on the morning of demo day.

This layer *needs* the Tier-1 + Tier-2 real-data rows above to anchor against. Specifically:

- Synthetic `BreedingRecommendation` rows pick species + populations from whatever real `ExSituPopulation` rows exist. No populations → no realistic targets → empty Panel 6.
- Synthetic `Transfer` rows pick source + destination from whatever `Institution` rows exist. One institution → no transfers possible.
- Synthetic `BreedingEvent` rows attach to populations. No populations → empty Panel 7.

Tier 1 alone (2 institutions, 2 populations) lets the synthetic layer seed a minimal coherent story. Tier 2 (5-10 populations across 3-5 institutions) gives it noticeably more variety. Tier 3 is gravy.

The synthetic layer is **safe** to run repeatedly — it's idempotent and rolls back cleanly via `seed_demo_coordination --clear`. Real operator-entered rows are never matched, never modified, never deleted.

---

4. Panel-by-panel expectations + demo lines

What each panel will look like at ABQ given Tier-1 + Tier-2 entry, and what to say.

Panel 1 — Coverage gap

What it shows: `~70 of ~75 threatened species have no known captive population.` (Exact number depends on what you enter.)

Demo line: *"This is the headline. Of every threatened endemic freshwater fish in Madagascar, we have ex-situ coverage on a small handful. The platform's first job is making this number visible — because the institutions that could close the gap don't currently share a registry."*

Panel 2 — Studbook status

What it shows: lopsided buckets — `Studbook-managed: 2` (the EAZA EEPs), `Breeding, not studbook: ~5` (CARES cichlids), `Holdings only: ~3`, `No captive population: ~70`.

Demo line: *"This isn't a uniform picture. Two species are under formal EAZA studbook management. A handful sit with hobbyist breeders under the CARES network. Most have no captive population at all. The platform makes that distribution legible — which is the prerequisite for any coordinator deciding where to direct effort."*

Panel 3 — Sex-ratio risk

What it shows: maybe 1-2 populations flagged, depending on whether you enter `count_male` / `count_female` counts.

Demo line: if anything is flagged: *"For populations where we have demographic detail, the platform flags ratio risk. Most of our records don't have that detail yet — that's an ask of the keeper community."* If nothing is flagged, skip the panel.

Panel 4 — Stale census

What it shows: likely empty if you enter `last_census_date = today` on every Tier-1 + Tier-2 row.

Demo line: skip the panel, or: *"This panel shows populations whose census is overdue. Today everything is fresh because we just entered it. The point is the platform will warn coordinators when records drift — that's*

how you keep a registry honest over time."

Panel 5 — Transfer activity

What it shows: `seed_demo_coordination` provides 2-4 synthetic in-flight transfers between the institutions you entered. All tagged `[DEMO_SEED]`.

Demo line: *"For the demo we've seeded a few synthetic transfers to illustrate the lifecycle — proposed, approved, in transit, completed. In production, coordinators log these as movements happen."*

Panel 6 — Open breeding recommendations

What it shows: 3-4 synthetic recommendations across the priority levels. Tagged `[DEMO_SEED]`.

Demo line: *"This is the coordinator's to-do list. Each recommendation has a priority, a target institution, and a due date. Critical recommendations surface first; overdue ones are tagged. This is the operational core of the dashboard."*

Panel 7 — Recent reproductive activity

What it shows: 5-10 synthetic events — spawnings, hatchings, mortalities — across the populations you entered.

Demo line: *"The event ledger. Spawnings and mortalities are logged here as they happen, with signed count deltas. Over time this becomes the longitudinal data layer underneath everything else."*

---

5. The honest-vs-padded decision

A judgment call, not a binary.

- **Honest** — enter Tier 1, optionally Tier 2, no Tier 3. Run the synthetic layer. Demo says "this is what we can see, much of which is empty, and that's the point." Strongest narrative for a conservation-funder audience. Risk: panels look sparse on their own.
- **Padded** — enter Tier 1 + 2 + 3. Demo looks fuller. Risk: someone asks where the Tier-3 data came from, and "I added it from public sources to round out the demo" lands awkwardly compared with "everything you see is sourced."

Recommendation: **honest**. The audience at ABQ is sophisticated enough that empty buckets read as *the problem the platform solves*, not *the platform doesn't have data yet*. Make sure the demo line for Panel 2 explicitly names the gap as the point — that flips the read.

---

6. Order of operations before ABQ

1. **Now → 2026-05-15.** Enter Tier 1 (alongside §2.7 EAZA EEPs) + whatever Tier 2 you have CSV-quality data for. ~1-2 hours total.
2. **Week of 2026-05-25.** Run §4.1 dry-run from the `WORKSHOP_DEMO_RUNBOOK`. Decide on Panel 4 (skip vs. include) and Panel 3 (skip vs. include) based on what's actually there.
3. **2026-05-31 (one day before).** Run `seed_demo_coordination` on staging per the morning-of checklist. Verify all panels render.
4. **2026-06-01 → 2026-06-05.** Demo with confidence in the data you have, named for what it is.